

Web Data Collection with R

Regular Expressions / RegEx

Peter Meißner / 2016-02-29 – 2016-03-04 / ECPR WSMT

How Regular Expressions work . . .

Patterns

Functions

Character Encodings

How Regular Expressions work . . .

What is it all about?

1. Regular Expressions refer to combination of two things

- ▶ a **syntax** that allows to define string patterns
 - ▶ e.g.: "[pP]eter", "\\d{4}-\\d{1,2}-\\d{1,2}"
- ▶ a set of **functions** doing string handling
 - ▶ base R has `grep()`, `grepl()`, `substring()`, ... nice because of options `ignore.case` and `invert` and because build in
 - ▶ more convenient stringr/stringi functions: `str_detect()`, `str_replace()`, `str_extract()`, ...

Patterns

Patterns

2. Regular Expressions providing string patterns

pattern	description
"Hallo"	1:1
"."	any character
"[]"	placeholder for one character
"[abc]"	set of characters (e.g. a,b, and c)
"[a-z]"	range of characters (e.g. a-z, not è, ä, ...)
"a*" / "a+"	none or more / one or more
"a{2,4}"	two up to four
"ac b"	ac or b
"[^ab]"	non of those
"^a"	starting with a
"a\$"	ending with a

Special Characters

3. Expressing Patterns

pattern	description
"\n"	newline
"\r"	carriage return
"\t"	tab
"\b"	word boundary (between "\\w" and "\\W")
"\122"	[matches ASCII character number 82 (octal)
"\x52"	matches ASCII character number 82 (hexadecimal)
"\u0052"	matches Unicode character number 82 (hexadecimal)

Character Classes

3. Expressing Character Classes

pattern	description
"\\d" / "\\D"	digit / no digit
"\\w" / "\\W"	word char. / no word char
"\\s" / "\\S"	white space char. / no ws char
"\\p{Currency_Symbol}"	unicode groups and blocks
"[:digit:]"	digit
"[:alpha:]"	characters (also è)
"[:alphanum:]"	word char.
"[^:alphanum:]"	white space char.

Syntax Characters

4. some characters have special meaning and cannot be used literally

► . \$ ^ { [(|)] } * + ?

character	description	matching
"\"	escapes "\", extra chars	<code>grep("\\\\\", "\\")</code>
"{"	numeral classifier	<code>grep("\\{", "{")</code>
...

Functions

functions for string detection

5. base functions for string detection / manipulation

name	description
<code>grep()</code>	searches for pat. and returns numeric index/content
<code>grep1()</code>	searches for pat. and returns logical index
<code>gregexpr()</code>	gives back each position of match
<code>nchar()</code>	length of string
<code>substr()</code>	extracts sequence of characters
<code>sub()</code>	replace one pat. match in string with other string
<code>gsub()</code>	replace all pat. matches in string with other string
<code>paste()</code>	concatonates vector elements into one string
<code>paste0()</code>	concatonates vector elements into one string
-	duplicates string
-	removes leading /trailing whitespace
-	adds whitespace to left, right, or both
-	returns matrix of strings x matches + 1
<code>cat()</code>	prints text to screen

functions for string detection

6. stringr/stringi functions for string detection / manipulation

name	description
-	searches for pat. and ret. numeric index/cont.
str_detect()	searches for pat. and returns logical index
str_locate()	gives back each position of match
str_length()	length of str.
str_sub()	extracts sequence of characters
str_replace	repl. one pat. match in str. with other str.
str_replace_all()	repl. all pat. matches in str. with other str.
-	concatonates vector elements into one str.
str_c()	concatonates vector elements into one str.
str_dup	duplicates string
str_trim	removes leading /trailing whitespace
str_pad	adds whitespace to left, right, or both
str_match	returns matrix of strings x matches + 1
cat()	prints text to screen

Character Encodings

Character Encodings

Character Encodings are ...

- ▶ are like family ...
- ▶ ... some of them you do not like but cannot avoid ...
- ▶ ... something we will struggle with but have cope anyways

The best thing is ...

- ▶ R has them all

The worst thing is ...

- ▶ R has them all

Character Encodings

- ▶ computers store everything as 0s and 1s (bits)
- ▶ in cs there are differing layers of abstraction
- ▶ one bit of information is called bit
- ▶ bits are quite uninformative as they only have two states
- ▶ so they are grouped into bytes (8 bits)
- ▶ one byte can have 256 different values (2^8)
- ▶ so it can store numbers 0 to 255 or 1 to 256 or ... -127 to 128
- ▶ or it can map to characters e.g. ASCII
(abcABC.:_-;#'+*~|<>i'\$%&/()=?)[{}^°, ...")
- ▶ ASCII is a character set - the set of characters you want to be able to store - even 7 Bits would suffice to store it

Character Encodings

- ▶ for larger character sets than ASCII (ä ö ü é è . . .) one needs to get clever since one byte does not suffice to map all characters to 0s and 1s
- ▶ unfortunate people got clever in differing ways
 1. using more than one byte to map more characters ('wide' characters, UTF-16, UCS-2, Windows OSs)
 2. using one or more bytes and using the first byte to encode how many are used ('multi-byte characters', UTF-8, Unix based OSs)
- ▶ otherwise we would not have to talk about character sets and character encodings

Character Encodings

```
rawToBits(as.raw(62:66)) # as bits
```

```
## [1] 00 01 01 01 01 01 01 00 00 01 01 01 01 01 01 00 00 00  
## [24] 00 01 00 00 00 00 00 00 01 00 00 01 00 00 00 00 01 00
```

```
as.raw(62:66) # bytes as hexa-decimal
```

```
## [1] 3e 3f 40 41 42
```

```
as.numeric(as.raw(62:66)) # as numbers
```

```
## [1] 62 63 64 65 66
```

```
rawToChar(as.raw(62:66)) # bytes as characters
```

```
## [1] ">?@AB"
```

A character set problem

```
text      <- rawToChar(as.raw(228))  
Encoding(text) <- "UTF-8"  
text
```

```
## [1] "\xe4"
```

```
Encoding(text) <- "latin1"  
text
```

```
## [1] "ä"
```

Results differ because for latin1 character 228 is known but not for UTF-8

An encoding problem

Of course UTF-8 knows how to encode “ä” ...

```
text <- "ä"  
charToRaw(text)
```

```
## [1] c3 a4
```

```
Encoding(text) <- "latin1"  
text
```

```
## [1] "Ãä"
```

... but here the results differ because “UTF-8” has another system translating characters to bytes. In latin1 the two bytes are interpreted as two characters.

Which default encoding does your R use

```
Sys.getlocale()
```

```
## [1] "LC_CTYPE=de_DE.UTF-8;LC_NUMERIC=C;LC_TIME=de_DE.UTF-8"
```

```
# if your locale is something other than UTF-8,  
# switch 'latin1' and 'UTF-8' and you shall be good to go
```

Changing interpretation of bytes

```
text <- "Små grodorna, små grodorna är lustiga att se."  
Encoding(text) <- "UTF-8"  
text
```

```
## [1] "Små grodorna, små grodorna är lustiga att se."
```

Changing interpretation of bytes

```
text <- "Små grodorna, små grodorna är lustiga att se."  
Encoding(text) <- "latin1"  
text
```

```
## [1] "SmÃ¥ grodorna, smÃ¥ grodorna Ãr lustiga att se."
```

Changing bytes and interpretation

```
text <- "Små grodorna, små grodorna är lustiga att se."  
text <- iconv(text, "UTF-8", "latin1")  
Encoding(text)
```

```
## [1] "latin1"
```

```
text
```

```
## [1] "Små grodorna, små grodorna är lustiga att se."
```

Noe that all sources might have another encoding than your R default locale!

```
text <- "Små grodorna, små grodorna är lustiga att se."  
text <- iconv(text, "UTF-8", "latin1")  
writeLines(text, "text_latin1.txt", useBytes = TRUE)  
text <- readLines("text_latin1.txt")  
Encoding(text)
```

```
## [1] "unknown"
```

```
text
```

```
## [1] "Sm\xe5 grodorna, sm\xe5 grodorna \xe4r lustiga att
```

```
Encoding(text) <- "latin1"  
text
```

```
## [1] "Små grodorna, små grodorna är lustiga att se."
```